

AOS Assignment 5: Key value store Report

- Murali Kammili

In my assignment, I have implemented both Least Recently Used (LRU) and Adaptive Replacement Cache (ARC) management algorithms. My key-value store cache is of variable size and can be tuned based on the given input. The cache size is based on the parameters I supply in the program and memory for the cache is allocated using `xmalloc()` implemented in my last assignment. We can observe the performance of my implementation using cache monitoring operations. By increasing the cache size, I could achieve zero '0' errors for both sets and gets. However, it may not be practically feasible to have a large cache size such that no single request fails for all types of input.

In the `kv_set()`, I am storing the key-value pair in the cache as soon as I get the request. First, I check if the cache already contains a value for the same key. Then we just update the new value of the respective key. If the key is already not present in cache, I try to place the new key-value pair in the cache. Eventually our cache gets filled and there will be no more space to accommodate the set request. In such scenario, we find a victim key-value pair that has to be evicted to save the new key-value pair. This eviction policy differs from the algorithm implemented. In my LRU program, I am evicting the least recently used page among the keys of same hash index. For ARC, I am maintaining an extra cache of the same size as my main cache. This extra cache is split into two equal lists (ghost caches) which are represented as `g1` and `g2`. These save the history of evicted pages from the main cache which is logically viewed as combination of `T1` and `T2`. This abstraction is represented as an integer value {1 for `T1`, 2 for `T2`} in my program using a variable '`T`' in '`xcache_k`' structure. I have declared a global variable '`P`', which allows us to maintain track of changes in `T1` and `T2`. If any key is found to be accessed more than once, we change it from `T1` to `T2` and `P` is changed accordingly.

In the `kv_get()`, I am first obtaining the hash index to narrow the key search in the cache. If I have found the key, then I return the value associated with that key. If it is a miss, I am incrementing the get error count and further continue to search in the ghost caches. If we could not find the key in the main cache, then I am scanning the ghost cache `g2`. If I found the key in `g2` then, I am evicting from `T1` to accommodate one more extra key-value pair for `T2`. Variable `P` maintains how many caches are reserved for `T2` at that moment. Similarly if I found the key in `g1` then, I am evicting from `T2` to accommodate one more extra key-value pair for `T1`. This eviction is again implemented using LRU.

In the `kv_delete()`, I am searching for the input key in cache and deleting the associated key-value pair by setting the values to null character. If I could not find the key, I am returning boolean value `False`.

In the `kv_reset()`, I am freeing all the allocated memory created using `xmalloc()` and resetting the values. In the end I am calling `xheap_snapshot()` which is implemented in the previous assignment to verify if all the allocated memory is freed up to be utilized again.

In `kv_init()`, I am calling `xmalloc_init()` first to initialize the buffer pool. I am allocating required cache memory by calling `xmalloc()`. Note that ARC requires double the cache size required by LRU.

In `get_cache_info()`, I am returning the required variable by pre-computing the values in the related functions inside the program. These capture the given cache monitoring operations and return the result.

The `most_popular_keys()`, returns the 'k' most frequently used keys after eliminating the duplicates.

I tested my both programs on the given test cases covering few scenarios and I found my ARC implementation to be a better performer than LRU. In my program the maximum size of my key is 64 bytes and maximum size of the value will be 1 KB as mentioned in the assignment. If the size is beyond that I have written condition to ignore it rather than truncation.

Sample :

For cache size - 25*91 and on input file xlarge

Using LRU :

```
total_hits 82
total_accesses 100
total_set_success 2048
cache_size 2226176
num_keys 2034
total_evictions 14
```

Set errors: 0, Get errors: 18

Using ARC :

```
total_hits 100
total_accesses 100
total_set_success 2048
cache_size 2211854
num_keys 2034
total_evictions 14
```

Set errors: 0, Get errors: 0

For cache size - 25*92 and on input file xlarge

Using LRU and ARC:

```
total_hits 100
total_accesses 100
total_set_success 2048
cache_size 2226176
num_keys 2039
total_evictions 9
```

Set errors: 0, Get errors: 0